

# boa



## Buffer Overrun Analyzer

Edo Cohen  
039374814  
sedoc@t2

Tzafrir Rehan  
039811880  
tzafrir@cs

Gai Shaked  
036567055  
gai@tx

March 2, 2011

# Chapter 1

## Introduction

### 1.1 Goal

Given a C program that performs buffer manipulations, statically identify whether the program may perform array access out of the array bounds.

# Chapter 2

## boa

### 2.1 Overview




Figure 2.1: Main components and stages

### 2.2 Constraint Generator

#### 2.2.1 Integers

#### 2.2.2 Direct array access

```
1  char buf[10];  
2  buf[10] = 'a';
```

#### 2.2.3 String manipulation functions

```
1  #include "string.h"  
2  
3  int main() {  
4      char *str1 = "longer_than_ten", *str2 = "short";  
5      char buf1[10], buf2[10];  
6      strcpy(buf1, str1);  
7      strcpy(buf2, str2);  
8  }
```

#### 2.2.4 Buffer aliasing

### 2.3 Linear Solver

The constraints generated represent a linear problem, and each solution of the problem suggests a set of ranges for the values each integer may receive and the allocation and usage of each buffer. As we aim to

find the tightest ranges, we direct our linear solver to find a solution maximizing -

$$Goal = \sum_{\text{Buffers}} \left[ \{buf!used!min\} + \{buf!alloc!min\} - \{buf!used!max\} - \{buf!alloc!max\} \right]$$

A solution satisfying this goal will maximize the lower bounds and minimize the upper bounds of each buffer access, and thus assure we get the tightest solution.

Once we have the solution<sup>1</sup> we test each buffer to verify that -

$$\{buf!used\} \subseteq \{buf!alloc\}$$

Which means -

$$\begin{aligned} \{buf!used!max\} &< \{buf!alloc!min\} \\ \{buf!used!min\} &\geq 0 \end{aligned}$$

Note that we stick to the size and numbering conventions of C, safe access to a buffer of size  $n$  is any access to the cells  $0 \dots n - 1$ . If the solution does not satisfy one of the constraints, we report a possible buffer overrun in this specific buffer.

### 2.3.1 Handling infeasible problems

In many cases, the constraints we generate create an infeasible linear problem. The simplest example of such case is -

```
1  int i;
2  i++;
```

The constraint generated from the second line will be -

$$\begin{aligned} \{tmp!max\} &\geq \{i!max\} + 1 \\ \{i!max\} &\geq \{tmp!max\} \end{aligned}$$

Which is obviously an infeasible set of constraints. The same problem holds in many different cases, including *strcat* (which concatenates one string to the end of another, and therefore implies an equivalent set of constraints to the string used length).

When our linear solver discovers that the constraints problem we have generated is infeasible, we wish make the smallest change to the problem and make it feasible once again. There is a great body of work in the area of finding and eliminating IIS (*irreducibly inconsistent system*), and we follow the algorithms and terminology of Chinneck and Dravnieks[3]. The common and naive approach is the deletion filtering -

1. input:  $Q$  is an infeasible set of constraints
2. Try to delete - for each  $q_i \in Q$  DO:
  - (a) Test whether  $Q \setminus q_i$  is feasible -
    - i. If infeasible - set  $Q = Q \setminus q_i$
3.  $Q$  is an IIS

After one iteration the algorithm return an IIS, which can be removed from the original problem. In case there are several IISs in the original problem - the algorithm should repeat until the problem become feasible. We have implemented this approach at first, and it did work well on small pieces of code, but naturally did not scale well - on the same testing system described in chapter 4 it took more than half an hour to eliminate the IISs in the 400 lines of source of *md5* library, and more than 8 hours to find the blames (section 2.3.2) as well. Therefore we read further and implemented an elastic filter for eliminating IIS. The main idea behind elastic filtering is adding a new *elastic variable* to each constraint, allowing it to *stretch* and therefore the infeasibility removed, than we solve the new problem, trying to minimize effect of the *elastic variables*

---

<sup>1</sup>The solution is a set of integer values, one for each of the problem variables, such that all the constraints are satisfied and the *Goal* value is maximized

1. Initialize  $S = \emptyset$  (will hold the IIS)
2.  $Q$  is an infeasible set of constraints of the form -

$$q_i : \sum_{j=1}^{n_i} a_{ij} X_{ij} \geq c_i$$

where  $a_{ij}, c_i$  are constants and  $X_{ij}$  is a variable of the constraint problem.

3. Add an *elastic variable*  $e_i$  to each constraint  $q_i \in Q$  such that -  $q_i : \sum_{j=1}^{n_i} a_{ij} X_{ij} + e_i \geq c_i$
4. Limit the elastic variables to accept non-negative values, and set the goal of the linear problem to minimize the sum of the elastic variables -  $Goal = - \sum e_i$
5. While the problem is feasible -
  - (a) Solve the linear problem, for each elastic variable  $e_i$  -
    - i. If  $e_i > 0$  -
      - A.  $S = S \cup \{q_i\}$
      - B. remove  $e_i$  ( $q_i : \sum_{j=1}^{n_i} a_{ij} X_{ij} \geq c_i$ )
6.  $S$  is an IIS

The elastic filtering algorithm based on

### 2.3.2 Blame system

## 2.4 Implementation

# Chapter 3

## Development Process

### 3.1 Initial Requirements

We set out to develop Boa by first defining a preliminary set of requirements:

- Provide analysis over any valid<sup>1</sup> C code as is, without requiring the programmer to provide any meta information.
- Provide *Soundness*: Report 100% of the buffer overruns in the code, with no *False Negatives* reported.

### 3.2 Lenient Assumptions

In order to achieve these requirements we also defined lenient assumption on the input code:

- The programmer knows that the C string library requires that a string must end with the NUL terminating character `'\0'`, and will never mutate the last byte of a buffer in a way that will cause an overrun.
- The programmer never uses an uninitialized value.

### 3.3 Research and Technology Survey

After defining the initial requirements we began researching past work on the subject of buffer overrun static analysis, and set out to find which tools already exist that can be used to build Boa.

Our research reached the conclusion of using a linear problem solver, solving a set of linear constraints generated from each instruction in the source code, in order to find buffer overruns, as described in chapter 2. TODO: Reference.

We set a goal of using only open source tools which are publically available for free use under an Open-Source compatible license, and converged towards the use of the following tools:

- GLPK - the GNU Linear Programming Kit [<http://www.gnu.org/software/glpk/glpk.html>], available under version 3 of the GNU Public License. Used to solve the linear constraints using the Simplex Algorithm.
- The Clang C Front-End [<http://clang.llvm.org/>], available under the University of Illinois/NCSA Open Source License. We planned on using Clang's plug-in system in order to go over the code's Abstract Syntax Tree and generate linear constraints according to the instructions in the code.

We then set out to create prototypes. First we created example, “*Hello World*” style, programs that make use of Clang and GLPK's interfaces separately,

---

<sup>1</sup>We define *Valid C* code as any code that is compiled by *gcc* with the *-Wall* flag without any warnings.

### 3.4 Test system

### 3.5 Version control

#### 3.5.1 Code reviews

# Chapter 4

## Results

We tested *boa* on several widespread real world programs. We tested to see whether *boa* discovers real buffer overruns, and also to evaluate the number of false alarms and their main causes. The source files used in all of the experiments are available in *boa* git repository[2].

Table 4.1 summarizes the performance of *boa* on the programs we present in this document, the reported running times are the results of experiments ran on a Dell Vostro 1310 laptop, with Intel Core2 Duo CPU T8100 2.10GHz and 2GB RAM running Debian GNU/Linux Wheezy (7.0.0), clang 2.9, llvm 2.9 and GLPK 4.43. On this humble configuration *boa* can analyze few thousands lines of code within seconds, thus the use of elastic filter did pay off and *boa* can be used to efficiently analyze any reasonable piece of C code.

	<i>fingerd</i>	<i>flex</i>	<i>syslog</i>
Source lines	230		332
Constraints	2894		1206
Running time	2.508s		1.304s
Buffers	34		15
Overruns reported	6		8
Real overruns	1		1

Table 4.1: *boa* performance on various real world examples

### 4.1 *fingerd*

We tested *boa* using *fingerd*, unix finger daemon. We altered the current source code to reflect the well known buffer overrun, used by the Internet worm in 1988. The overrun is caused by using the unsafe function *gets* to read data into the 1024<sup>1</sup> bytes buffer *line*. As far as we know, this is the only real buffer overrun in the 230 lines of source code.

Running on *fingerd* source, *boa* reported overruns on 6 out of the 34 buffers. Next we present *boa*'s blame for three of them, and analyze the reason for the reported overrun -

**line** is the only real overrun in *fingerd*

```
line tests/realworld/fingerd/fingerd.c:85
- unsafe function call gets [tests/realworld/fingerd/fingerd.c:121]
- unknown function call realhostname_sa [tests/realworld/fingerd/fingerd.c:128]
- memchr call might read beyond the buffer [tests/realworld/fingerd/fingerd.c:139]
[ ... 10 more lines ... ]
```

The overrun discovered by *boa*, and the real cause reported briefly. Note another result of *gets* - every other buffer access based on *line*'s length will be reported as an overrun.

---

<sup>1</sup>Back in 1988 *line* was 512 bytes, but it does not matter for the analysis.



**rhost** is a char buffer meant to hold the host name

```
rhost tests/realworld/fingerd/fingerd.c:86
- unknown function call realhostname_sa [tests/realworld/fingerd/fingerd.c:128]
- unknown function call realhostname_sa [tests/realworld/fingerd/fingerd.c:128]
- buffer alias with offset [tests/realworld/fingerd/fingerd.c:128]
- buffer alias with offset [tests/realworld/fingerd/fingerd.c:128]
```

This false alarm is caused by the use of *realhostname\_sa*, from *socket.h*. This false alarm could be avoided if boa would model *socket.h* functions, but even now the output let the user identify the cause immediatly and decide manually wether this call is safe or not.

**malloc** is a generic name for any buffer created by a malloc call, one can distinguish between two malloc calls by their source location (filename and line number)

```
malloc tests/realworld/fingerd/fingerd.c:141
- buffer alias [tests/realworld/fingerd/fingerd.c:149]
- buffer alias [tests/realworld/fingerd/fingerd.c:149]
- buffer alias [tests/realworld/fingerd/fingerd.c:149]
- buffer alias [tests/realworld/fingerd/fingerd.c:141]
- buffer alias [tests/realworld/fingerd/fingerd.c:149]
```

This blame might seem wierd at a first look, how comes buffer alias alone cause an overrun? But the solution appears quickly by looking at the source lines (141, 149) reffered by the blame -

```
141      if ((t = malloc(sizeof(line) + 1)) == NULL)
```

...

```
149      for (end = t; *end; end++)
150          if (*end == '\n' || *end == '\r')
151              *end = '_';
```

The programmer allocates a buffer large enough to include *line*, and then iterates through the array using the ++ operator on a pointer. Since boa is a flow-insensitive analyzing tool, we can not assure that the incremental pointer aliasing will be limited to the buffer size - and therefore boa reports a possible buffer overrun.

## 4.2 flex

# Bibliography

- [1] Buffer Overrun Detection using Linear Programming and Static Analysis
- [2] boa git repository - <https://github.com/tzafrir/boa/>
- [3] Locating Minimal Infeasible Constraint Sets in Linear Programs